

# Introduction to Python

Parham Alvani

Amirkabir University of Technology

`parham.alvani@gmail.com`

May 14, 2015

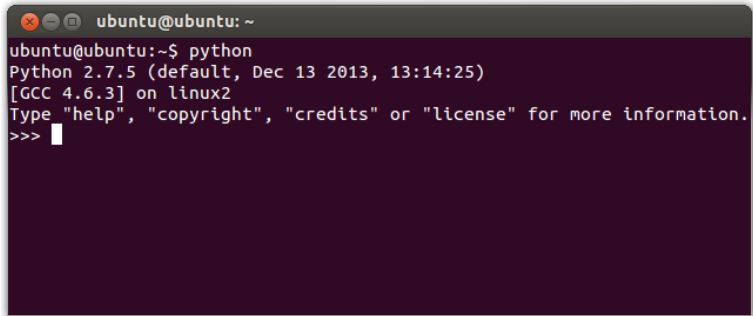


## what is python...

Python is a high-level, interpreted, interactive and object-oriented scripting language. Python is designed to be highly readable..



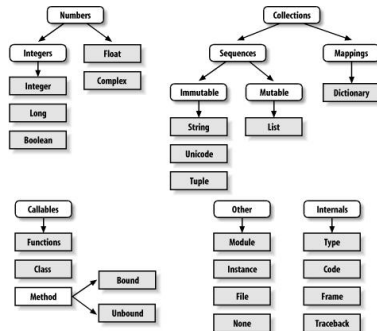
# Environment

A terminal window with a dark purple background and a grey title bar. The title bar contains three window control icons (close, minimize, maximize) and the text 'ubuntu@ubuntu: ~'. The terminal text is as follows:

```
ubuntu@ubuntu:~$ python
Python 2.7.5 (default, Dec 13 2013, 13:14:25)
[GCC 4.6.3] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> 
```

# Variable Types

- ▶ Numbers
- ▶ String
- ▶ List
- ▶ Tuple
- ▶ Dictionary



# Numbers

- ▶ Number data types store numeric values.
- ▶ They are immutable data types, means that changing the value of a number data type results in a newly allocated object.

# Strings

- ▶ Strings are amongst the most popular types in Python.
- ▶ We can create them simply by enclosing characters in quotes.
- ▶ Python treats single quotes the same as double quotes.

- ▶ The most basic data structure in Python is the sequence.
- ▶ Each element of a sequence is assigned a number - its position or index.
- ▶ The first index is zero, the second index is one, and so forth.

# Tuples

- ▶ A tuple is a sequence of immutable Python objects.
- ▶ Tuples are sequences, just like lists.
- ▶ The differences between tuples and lists are :
  - the tuples cannot be changed unlike lists
  - tuples use parentheses, whereas lists use square brackets.



# Dictionary

- ▶ Each key is separated from its value by a colon (:) )
- ▶ the items are separated by commas
- ▶ and the whole thing is enclosed in curly braces.

- ▶ If-Then-Else
- ▶ For
- ▶ While
- ▶ Exceptions

# Functions

- ▶ Function blocks begin with the keyword `def`, followed by the function name and parentheses.
- ▶ Any input parameters or arguments should be placed within these parentheses.
- ▶ The first statement of a function can be an optional statement - the documentation string of the function or docstring.

# Functions

```
def square(x):  
    return x * x
```


```
def square():  
    return "Hello"
```

- ▶ The class statement creates a new class definition.
- ▶ The name of the class immediately follows the keyword `class` followed by a colon as follows

..., but there are problems with relying on data too much.

Not everything that can be counted counts, and not everything that counts can be counted.

- Albert Einstein




However, any data is better than none.

An approximate answer to the right problem is worth a good deal more than an exact answer to an approximate problem.

- John Tukey



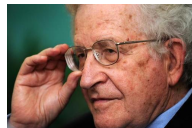
# Big Data, the Noam Chomsky Way



Big data is a step forward. But, our problems are not lack of access to data, but understanding them. [Big data] is very useful if I want to find out something without going to the library, but I have to understand it, and that's the problem.



# Big Data, the Noam Chomsky Way



Big data is a step forward. But, our problems are not lack of access to data, but understanding them. [Big data] is very useful if I want to find out something without going to the library, but I have to understand it, and that's the problem.

Hmmm, not very much Chomsky-ish ..., but wait!

# Big Data, the Noam Chomsky Way



Big data is a step forward. But, our problems are not lack of access to data, but understanding them. [Big data] is very useful if I want to find out something without going to the library, but I have to understand it, and that's the problem.

Hmmm, not very much Chomsky-ish ..., but wait!

We can be confident that any system of power - whether it's the state, Google, or whatever - is going to use the best available technology to control, to dominate, and to maximize their power. And they'll want to do it in secret.

Now that's sounding more like Chomsky.

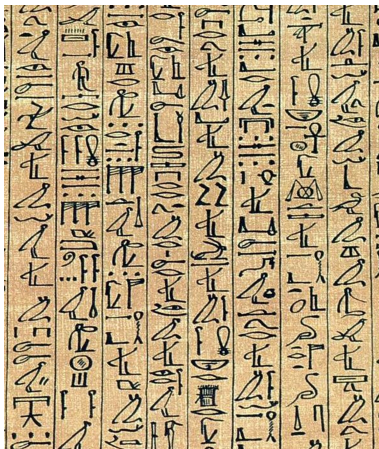
# They Want to Do It In Secret ...

The truth cannot stay hidden forever!



# A Brief History of Data Management!

- ▶ Manual recording
- ▶ From tablets to papyrus, to parchment, and then to paper




- ▶ Gutenberg's printing press



# 1800's - 1940's

- ▶ Punched cards (no fault-tolerance)
- ▶ Binary data
- ▶ 1890: US census
- ▶ 1911: IBM appeared



## 1940's - 1970's

- ▶ Magnetic tapes
- ▶ Batch transaction processing
- ▶ File-oriented record processing model (e.g., COBOL)
- ▶ Hierarchical DBMS (one-to-many)
- ▶ Network DBMS (many-to-many)





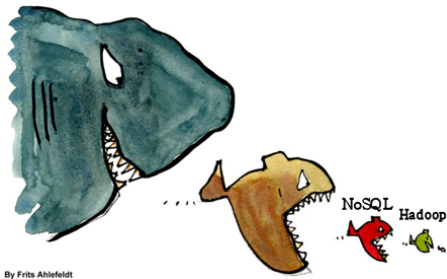
- ▶ Relational DBMS (tables) and SQL
- ▶ ACID
- ▶ Client-server computing
- ▶ Parallel processing



## ► The Internet...

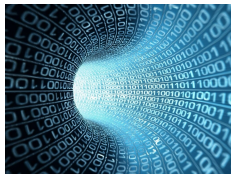


- ▶ NoSQL: BASE instead of ACID
- ▶ Big Data



# Big Data

- ▶ In recent years we have witnessed a **dramatic increase** in available data.
- ▶ For example, the **number of web pages** indexed by Google, which were around **one million** in 1998, have exceeded **one trillion** in 2008, and its expansion is accelerated by appearance of the social networks.




# Big Data Definition

- **Big Data** refers to datasets and flows **large enough** that has outpaced our capability to **store, process, analyze, and understand**.



# The Four Dimensions of Big Data

- ▶ **Volume**: data **size**
- ▶ **Velocity**: data **generation rate**
- ▶ **Variety**: data **heterogeneity**
- ▶ **Veracity**: **uncertainty** of accuracy and authenticity of data



## Big Data Market Driving Factors

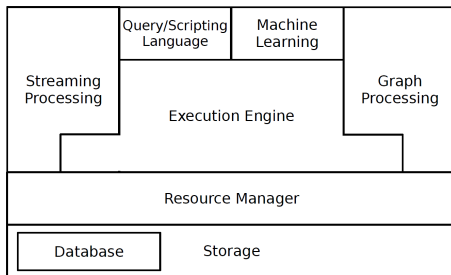
- ▶ Mobile devices
- ▶ Internet of Things (IoT)
- ▶ Cloud computing
- ▶ Open source initiatives



# The Big Data Stack!




# Big Data Analytics Stack



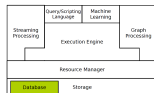
# Big Data - Storage (Filesystem)

- ▶ Traditional filesystems are not well-designed for large-scale data processing systems.
- ▶ **Efficiency** has a higher priority than other features, e.g., directory service.
- ▶ Massive size of data tends to store it across **multiple machines** in a distributed way.
- ▶ HDFS, Amazon S3, ...



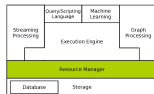
# Big Data - Database

- ▶ Relational Databases Management Systems (RDMS) were **not** designed to be distributed.
- ▶ **NoSQL** databases **relax** one or more of the **ACID** properties: **BASE**
- ▶ Different data models: **key/value**, **column-family**, **graph**, **document**.
- ▶ Dynamo, Scalaris, BigTable, Hbase, Cassandra, MongoDB, Volde-mort, Riak, Neo4J, ...



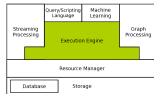
# Big Data - Resource Management

- ▶ Different frameworks require different **computing resources**.
- ▶ Large organizations need the ability to **share data and resources** between multiple frameworks.
- ▶ **Resource management** share resources in a cluster between **multiple frameworks** while providing resource **isolation**.
- ▶ Mesos, YARN, Quincy, ...



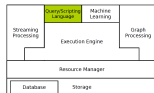
# Big Data - Execution Engine

- ▶ **Scalable** and **fault tolerance** parallel data processing on clusters of unreliable machines.
- ▶ Data-parallel **programming model** for clusters of commodity machines.
- ▶ MapReduce, Spark, Stratosphere, Dryad, Hyracks, ...



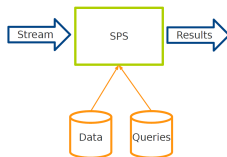
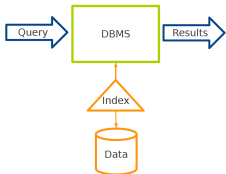
# Big Data - Query/Scripting Language

- ▶ **Low-level** programming of execution engines, e.g., MapReduce, is **not** easy for end users.
- ▶ Need **high-level** language to improve the query capabilities of execution engines.
- ▶ It translates **user-defined** functions to **low-level** API of the execution engines.
- ▶ Pig, Hive, Shark, Meteor, DryadLINQ, SCOPE, ...

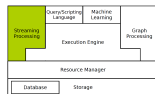


# Big Data - Stream Processing

- ▶ Providing users with **fresh** and **low latency** results.
- ▶ Database Management Systems (**DBMS**) vs. Stream Processing Systems (**SPS**)




- ▶ Storm, S4, SEEP, D-Stream, Naiad, ...



# Big Data - Graph Processing

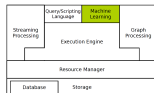
- ▶ Many problems are expressed using **graphs**: sparse **computational dependencies**, and **multiple iterations** to converge.
- ▶ Data-parallel frameworks, such as MapReduce, are not ideal for these problems: **slow**
- ▶ Graph processing frameworks are **optimized** for graph-based problems.
- ▶ Pregel, Giraph, GraphX, GraphLab, PowerGraph, GraphChi, ...



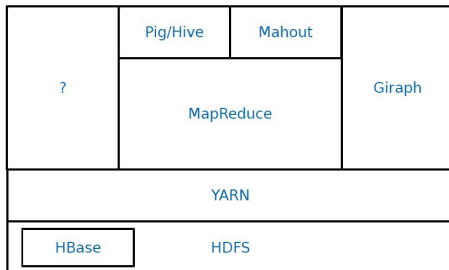


# Big Data - Machine Learning

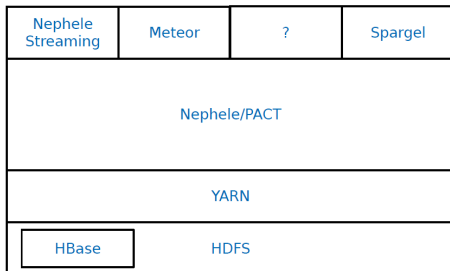
- ▶ Implementing and consuming machine learning techniques at scale are **difficult tasks** for developers and end users.
- ▶ There exist platforms that address it by providing scalable machine-learning and data mining libraries.
- ▶ Mahout, MLBase, SystemML, Ricardo, Presto, ...



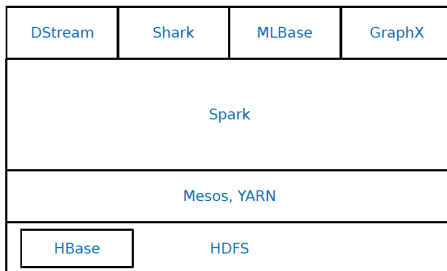
# Hadoop Big Data Analytics Stack



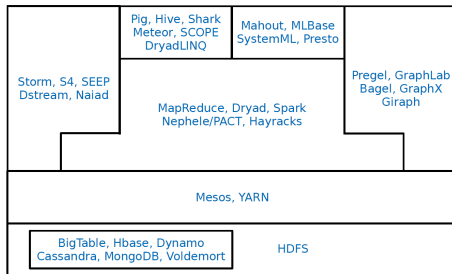
# Stratosphere Big Data Analytics Stack



# Spark Big Data Analytics Stack



# Summary



# Questions?